


Programming Languages

CSE 3302, Summer 2019
Module 02
Regular Expressions and
Context-Free Grammars



M02

*Regular Expressions and
Context-Free Grammars*



Formality of Programming Languages

- Unlike natural languages, programming languages must be *precise* and without *ambiguity*.
 - Given a source program, the exact steps to execute it must be determinable in a predictable, repeatable fashion.
 - Otherwise how would the computer know what to do and how would the programmer know what the computer will do?
- Formal notation is therefore used to define the language and then to write programs in that language.
 - The notation describes the *syntax* and the *semantics*.



[“*Formal Notation*” ...]

- *Notation* is a way to express information, usually in a *written* form.
 - Using the Roman letters (A ... Z) to express the written form of English words is an example of a *notation*.
- *Formal* (here at least) means that the *notation* is sufficiently *precise* and *well defined* that it can be processed *mechanically*.
 - The symbols +, -, *, and / along with their accepted meanings are a “formal notation” for expressing basic arithmetic.
 - It’s not hard to write a program that accepts formulae in this notation and returns the corresponding values. The formulae are processed *mechanically*. (That is, no *human interpretation* is required.)



Programming Language Syntax

- *Syntax* is the *form* of the language as opposed to its *meaning* (the *semantics*)
- Why differentiate?
 - Languages tend to differ *greatly* in syntax but often have *very similar* aspects of semantics.
 - Easier to learn new language if the meaning behind the form can be distinguished.
 - Syntactic analysis has been extensively studied. There are many algorithms for determining the form of a program.

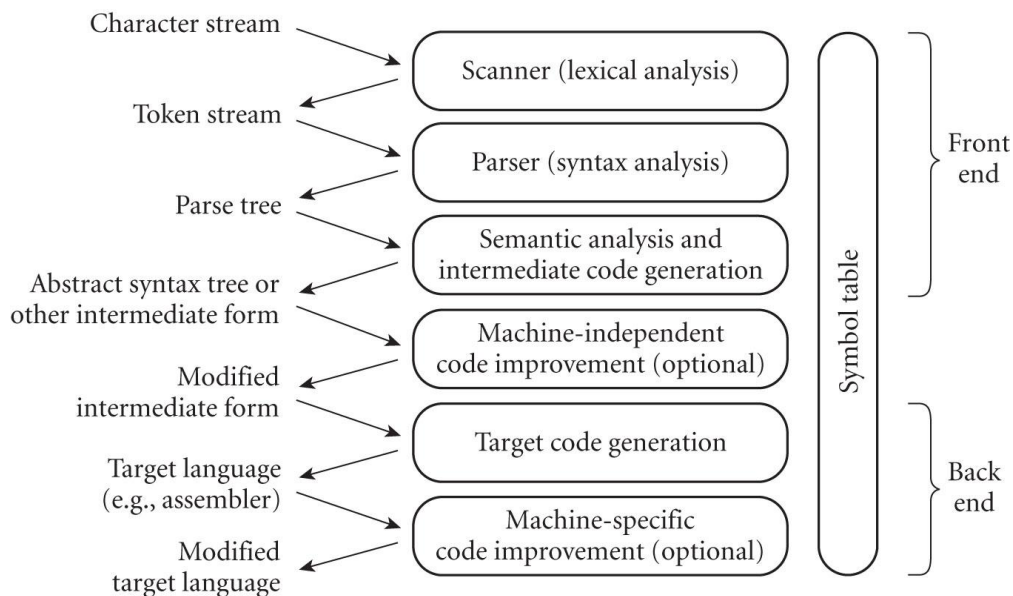


Specifying Syntax

- As seen in the compilation overview, the first two phases of compilation are the *lexical analysis* and the *syntactical analysis*.
 - Lexical Analysis converts characters into *tokens*.
 - Syntactical Analysis converts tokens into a *parse tree*.
- *Regular expressions* (RE) are used to specify the *token* formats.
- A *Context-Free Grammar* (CFG) is used to specify the parse structure.



Phases of Compilation



[*Grammars and the Chomsky Language Hierarchy*]

- A *grammar* specification includes

- A set of *terminal* symbols
- A set of *non-terminal* symbols
- A set of *production rules*
- A *start symbol*


Terminals	a, b
Non-terminals	S
Production rules	$S \rightarrow aSb$ $S \rightarrow \epsilon$
Start symbol	S

- Restrictions on the form of the production rules can be used to categorize grammars into classes, each of which is more *expressive* than the previous.
- The example expresses the language $a^n b^n$, for $n \geq 0$.



[*Grammars and the Chomsky Language Hierarchy*]

- The hierarchy shows four levels of grammar expressivity.
- Each properly includes the previous.
- The levels have increasing expressiveness, but also parsing complexity, resource requirements, etc.



[*Grammars and the Chomsky Language Hierarchy*]

Type	Name	Allowable Productions	Example Language	Example Grammar	Example Use	Recognizing Automaton	Storage Required	Parsing Complexity
0	Recursively Enumerable	Unrestricted				Turing Machine	Infinite Tape	Undecidable
1	Context Sensitive	$\alpha \rightarrow \beta$ where $ \alpha \leq \beta $ $\alpha \in V^*NV^*$ $\beta \in V^+$	$a^n b^n c^n$	$S \rightarrow aSBC$ $S \rightarrow aBC$ $CB \rightarrow BC$ $aB \rightarrow ab$ $bB \rightarrow bb$ $bC \rightarrow bc$ $cC \rightarrow cc$		Linear Bounded Automaton	Tape a linear multiple of input length	NP Complete
2	Context Free	$A \rightarrow \alpha$ $A \in N$ $\alpha \in V^*$	$a^n b^n$	$S \rightarrow aSb$ $S \rightarrow ab$	Arithmetic Expression $x = a + b * c$	Pushdown Automaton	Pushdown Stack	$O(n^3)$
3	Regular Right Linear Finite Automaton Recognizable	$A \rightarrow xB$ $A \rightarrow x$ $A, B \in N$ $x \in T^*$	$a^n b$	aa^*b or $S \rightarrow ab$ $S \rightarrow aS$	Identifier VECTOR7	Finite Automaton	Finite Storage	$O(n)$




[“Recursive” Production Rules ...]

- Later we’ll learn that *recursion* is what separates *Context-Free Grammars* from *Regular Expressions*. No recursion in REs.
- On the previous chart, $S \rightarrow aS$ appears as a rule for an RE.
 - Hey, isn’t that *recursive*? After all, S refers to itself, right?
- **No!** It just *looks* as if it’s recursive. :)
- Because the S appears on the far right (there’s nothing after S in the rule), this is just a way of expressing a Kleene-* operation (a^*).
- A proof that this isn’t true recursion requires going deeper in formal language theory than required for this class, so just accept it as true.



[John Backus, Peter Naur, and Noam Chomsky]



- ① <http://www.columbia.edu/cu/computinghistory/backus.html>
② <http://homepages.cs.ncl.ac.uk/brian.randell/NATO/N1968/NAUR.htm>
③ [https://commons.wikimedia.org/wiki/File:Noam_Chomsky_\(1977\).jpg](https://commons.wikimedia.org/wiki/File:Noam_Chomsky_(1977).jpg)



Tokens

- A *token* is the basic building block of a program.
 - The shortest strings of characters in the source program that have individual meaning.
- Tokens come in various types, according to the programming language's specification.
 - Common types include *numbers, identifiers, keywords, operators, punctuation marks*, and so forth.
- The formats of the tokens of a programming language are normally specified using *regular expressions*.



Simple Regular Expression

natural_number \rightarrow *non_zero_digit digit**

non_zero_digit \rightarrow 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9


digit \rightarrow 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9

- Describes any string of decimal digits one or more digits in length that does not start with a 0.



Definition of a Regular Expression (RE)

1. A *character*
2. The *empty string*, denoted as ϵ
3. The *concatenation* of two REs
Meaning one RE after the other RE
4. The *alternation* of two REs, denoted by $|$
Meaning one RE *or* the other RE
5. An RE followed by the *Kleene star*, denoted as $*$
Meaning zero or more repetitions of the RE



Stephen C. Kleene
/'klemi/ KLAY-nee

<https://www.nap.edu/read/9649/chapter/10>



Definition ...

- The set of strings defined by an RE is known as its *language*.
 - Not the same sense as in *programming language*!
- An RE is *not* allowed to *recurse*.
 - That is, no part of an RE can refer to itself, even indirectly.
 - If it did, it's no longer an RE, but a CFG instead.
 - (Remember the *Chomsky Language Hierarchy*.)
- There are some notational conveniences that don't change the expressive power of REs, but are very useful.



Notational Conveniences ...

- Kleene Plus is an example of a notational convenience.
 - It means *one or more* occurrences of a RE.
- It's easily represented using Kleene Star (*zero or more*).
 - If A is an RE, A^+ is defined to be AA^* .
- By expressing Kleene Plus in terms of Kleene Star, we see that we didn't gain any expressive power, just convenience.



Notational Conveniences ...

- There are other conveniences.
 - $A^? \equiv A|\epsilon$ or 0 or 1 occurrence of an RE.
 - $[abc] \equiv a|b|c$ or any of a set of characters.
 - (and there are others ...)
- Each of these conveniences is representable using the basic five ways of writing a regular expression, so no expressive power is gained by using them.



Extensions ...

- The definition of RE is often truly *extended* by string-processing programming languages.
 - Perl, Python, Ruby, awk, sed, ... all add extensions that increase the expressive power beyond the *regular sets*.
 - For example, referring back to an earlier matched string is not possible in a formal RE.
- Though these languages call their representations Regular Expressions, they are *not* true REs.



Numeric Literal Regular Expression

$number \rightarrow integer \mid real$

$integer \rightarrow digit\ digit^*$

$real \rightarrow integer\ exponent \mid decimal\ (exponent \mid \epsilon)$

$decimal \rightarrow digit^*\ (\cdot\ digit \mid digit\ \cdot\)\ digit^*$

$exponent \rightarrow (e \mid E) (+ \mid - \mid \epsilon)\ integer$

$digit \rightarrow 0 \mid 1 \mid 2 \mid 3 \mid 4 \mid 5 \mid 6 \mid 7 \mid 8 \mid 9$

- Q: Why the weird-looking definition for *decimal*?



Numeric Literal Examples

- *integer*
 - Simple! Just any non-zero length string of decimal digits.
 - 0, 1, 2, 3, 4567, 007, 070, ...
 - But **not**, e.g., -1 (there's no sign on a numeric literal).
- *real*
 - Trickier! An *integer* with an *exponent* or a *decimal* with an optional *exponent*.
 - *exponent* is e or E followed by an optionally signed *integer*.
 - *decimal* is a decimal point with digits on *at least* one side.
 - 0E+1, 2e-3, 4E5, 6., .7, 8.9E0, ...
 - But **not**, e.g., .E1 (need at least one digit next to the .)



Regular Expression Limitations

- Regular Expressions are *great*!
 - Simple to compose.
 - Simple to analyze.
 - Simple to decide if a given string matches a given RE.
 - *Linear* in the length of the string.
- *But*, they can't be used to check, e.g., *matching* or *nested* parentheses, brackets, etc.
- Why not?



Regular Expression Limitations

- A regular expression is trivially convertible to a non-deterministic finite automata (NFA) which is trivially convertible to a deterministic finite automata (DFA).
- The important word in those names is *Finite*.
 - The automata have a fixed, finite number of states.
 - No matter how big that number is, it's possible to construct a string that is *just a little bit more complicated* so the finite automata cannot recognize it.



Context-Free Grammar (CFG)

- Regular Expressions are limited in expressiveness.
 - Cannot define strings that are required to have arbitrary ...
 - Well-formed and/or nested parentheses, brackets, etc.
e.g., $([[\{()\}\{\}]]\{\})$
 - Matching pairs
e.g., $a^n b^n \rightarrow ab, aabb, aaabbb, \dots$
- Next step up is a *Context-Free Grammar*
 - Definitions can refer to themselves, i.e., can *recurse*.
 - We generally express CFGs in *Backus-Naur Form*.



[*Backus - Naur Form*]

- A *formal notation* for writing Context-Free Grammars.
- Originally developed to specify Algol and first used in the Algol 60 report.
- Looked a bit different from what we use here, but the expressive power is the same. We use ...
 - “ \rightarrow ” instead of “ $::=$ ”
 - *italics* for *non-terminals* instead of enclosing in $\langle \rangle$
 - ϵ for empty string
 - Won't bother with "" unless required for clarity

BNF

```
<postal-address> ::= <name-part> <street-address> <zip-part>
<name-part> ::= <personal-part> <last-name> <opt-suffix-part> <EOL>
| <personal-part> <name-part>
<personal-part> ::= <initial> "." | <first-name>
<street-address> ::= <house-num> <street-name> <opt-apt-num> <EOL>
<zip-part> ::= <town-name> ", " <state-code> <ZIP-code> <EOL>
<opt-suffix-part> ::= "Sr." | "Jr." | <roman-numeral> | ""
<opt-apt-num> ::= <apt-num> | ""
```

https://en.wikipedia.org/wiki/Backus-Naur_form

Ours

```
postal-address  $\rightarrow$  name-part street-address zip-part
name-part  $\rightarrow$  personal-part last-name opt-suffix-part EOL
| personal-part name-part
personal-part  $\rightarrow$  initial . | first-name
street-address  $\rightarrow$  house-num street-name opt-apt-num EOL
zip-part  $\rightarrow$  town-name , state-code ZIP-code EOL
opt-suffix-part  $\rightarrow$  Sr. | Jr. | roman-numeral |  $\epsilon$ 
opt-apt-num  $\rightarrow$  apt-num |  $\epsilon$ 
```



CFG Example

$expr \rightarrow id \mid number \mid - expr \mid (expr) \mid expr \ op \ expr$
 $id \rightarrow (_ \mid a \mid b \mid \dots \mid z)(_ \mid a \mid b \mid \dots \mid z \mid 0 \mid 1 \mid \dots \mid 9)^*$
 $op \rightarrow + \mid - \mid * \mid /$

- Notice that *expr* refers to itself. This definition is *recursive*.
 - And *expr* is *not* on the far right, so *true* recursion.
- Q: What set of strings does *id* define?



CFG Example

- To *derive* (or *generate*) a string from a CFG, begin with the start symbol and replace non-terminals according to the rules until only terminals remain.
 - A *sentential form* is the start symbol or any form derived from it.
 - A *sentence* is a sentential form which has only terminal symbols.
- Example, generate
 $\text{slope} * x + \text{intercept}$
using the *expr* CFG.

expr

expr *op* expr

expr op *id*

expr + *id*

expr *op* expr + *id*

expr op *id* + *id*

expr * *id* + *id*

id * id + id

slope * x + intercept



CFG Example

- That derivation was *Right-Most*.
 - We replaced the *right-most* non-terminal each time we took a step in the derivation.
 - *Except* for the final replacement of the *id* non-terminals with their actual words. (This was to make the derivation a little shorter and easier to fit on the slide.)
- A *Left-Most* derivation replaces the *left-most* non-terminal each time a step is taken in the derivation. (Duh.)




CFG Example

- The steps we took in the derivation correspond to the construction of a *parse tree*.
- The *root* of the parse tree is the *start symbol*.
- Every time we use a production rule, it's the same as adding a new (set of) node(s) to the tree.
- All internal nodes of the parse tree are *non-terminals*.
- The leaves of the final parse tree are *terminals*.
 - These terminals are the *tokens* of the original string.



CFG Example Parse Tree (from the Right)



CFG Example

- There's another way to do the generation of
slope * x + intercept
using the *expr* CFG.
- This derivation goes from the *left*
instead of the *right*.

expr

expr op expr

id op expr

id * expr

id * expr op expr

id * id op expr


id * id + expr

id * id + id

slope * x + intercept



CFG Example Parse Tree (from the Left)



CFG Example

- The CFG allows two parse trees for the same string because its definition is *ambiguous*.
 - Also, it allows incorrect parsing of operator precedence.
- So is the answer to always derive from the right?
 - After all, that got the correct derivation.
- No!
 - Consider the right-most derivation of
intercept + slope * x

expr

expr op expr

expr op id

expr * id

expr op expr * id

expr op id * id


expr + id * id

id + id * id

intercept + slope * x



CFG Example 2 Parse Tree (from the Right)



Improved CFG Example

- We need a CFG that is not ambiguous *and* honors our concept of operator precedence.
- A better (though more complex) CFG would be

$expr \rightarrow term \mid expr \text{ add_op } term$

$term \rightarrow factor \mid term \text{ mult_op } factor$

$factor \rightarrow id \mid number \mid - factor \mid (expr)$

$add_op \rightarrow + \mid -$

$mul_op \rightarrow * \mid /$

expr

expr add_op term

expr add_op factor

expr add_op id

expr + id

term + id

term mult_op factor + id

term mult_op id + id

term * id + id


factor * id + id

id * id + id

slope * x + intercept



Improved CFG Parse Tree (from the Right)



Improved CFG Example

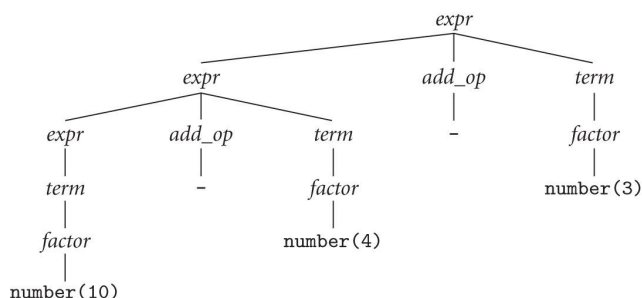

- What about doing the derivation from the left?
- We still get the correct derivation!
 - The grammar is unambiguous and it honors our intuitive understanding of operator precedence.
 - Multiply and divide are “tighter” than addition and subtraction.
- But what about *associativity*?
 - How are “like” operations grouped?

expr
expr add_op term
term add_op term
term mult_op factor add_op term
factor mult_op factor add_op term
id mult_op factor add_op term
id * factor add_op term
id * id add_op term
id * id + term
id * id + factor
id * id + id
slope * x + intercept



Improved CFG Example

- This CFG handles *precedence* and *associativity* according to our intuitive expectations.
 - 3+4*5 means
 - 3+(4*5) = 23
 - not (3+4)*5 = 35.
 - 10-4-3 means
 - (10-4)-3 = 3
 - not 10-(4-3) = 9.



Improved CFG Example

- OK, so it works. But *why* does it work?
- Operator *precedence* is enforced by having nested rules that permit the repetition of only certain operators at each precedence level.
 - *expr* is “looser” than *term* as it occurs higher in the CFG definition.
 - Therefore the *add_op* operators are “looser” than the *mul_op* operators.
- Operator *associativity* is enforced by having the *expr* and *term* rules recurse on their left side rather than the right.
 - Therefore operations group from the left instead of the right.

$expr \rightarrow term \mid expr \text{ add_op } term$

$term \rightarrow factor \mid term \text{ mul_op } factor$

$factor \rightarrow id \mid number \mid - factor \mid (expr)$

$add_op \rightarrow + \mid -$

$mul_op \rightarrow * \mid /$



Improved CFG

- In summary, we have seen that operator *precedence* and *associativity* are defined by how the CFG is constructed.
 - They are both *syntactic* properties of the language.
- *Precedence* is in *levels*.
 - Several operators may be at the same precedence level.
 - E.g., In C, + and -, * and /. (The latter are at a higher level.)
- Associativity is *left-to-right*, *right-to-left*, or *does-not-apply*.
 - E.g., in C, +, -, *, and / are L → R but +=, -=, *=, /= are R → L.
 - Any non-associative operators in C? What about Python?



Review Questions (1)

1. What is the difference between *syntax* and *semantics*?
2. What are the three basic operations that can be used to build complex *regular expressions* from simpler regular expressions?
3. What additional operation is provided in *context-free grammars*?
4. What is *Backus-Naur Form*? When and why was it developed?
5. Name a language in which indentation affects program syntax?



Review Questions (2)

6. When discussing context-free languages, what is a *derivation*? What is a *sentential form*?
7. What is the difference between a *right-most* derivation and a *left-most* derivation?
8. What does it mean for a CFG to be *ambiguous*?
9. What are *associativity* and *precedence*? Why are they significant in parse trees?

